

Behavior Knowledge Merge in Reinforced Agentic Models

Anonymous ACL submission

Abstract

Reinforcement learning (RL) is central to post-training, particularly for agentic models that require specialized reasoning behaviors. In this setting, model merging offers a practical mechanism for integrating multiple RL-trained agents from different tasks into a single generalist model. However, existing merging methods are designed for supervised fine-tuning (SFT), and they are suboptimal to preserve task-specific capabilities on RL-trained agentic models. The root is a task-vector mismatch between RL and SFT: on-policy RL induces task vectors that are highly sparse and heterogeneous, whereas SFT-style merging implicitly assumes dense and globally comparable task vectors. When standard global averaging is applied under this mismatch, RL's non-overlapping task vectors that encode critical task-specific behaviors are reduced and parameter updates are diluted. To address this issue, we propose **Reinforced Agent Merging (RAM)**, a distribution-aware merging framework explicitly designed for RL-trained agentic models. RAM disentangles shared and task-specific unique parameter updates, averaging shared components while selectively preserving and rescaling unique ones to counteract parameter update dilution. Experiments across multiple agent domains and model architectures demonstrate that RAM consistently retains specialized agent capabilities and outperforms existing baselines, achieving state-of-the-art performance. Code is available¹.

1 Introduction

Post-training has become a cornerstone for aligning large language models (LLMs) to diverse domains (Wei et al., 2021; Ouyang et al., 2022; Jaech et al., 2024). While specializing a model for a single task is effective, real-world applications typically require a single model to possess multi-task




Figure 1: Performance comparison of RAM/RAM+ and baselines on 12 tasks across three agent domains. Our method achieves the best average performance and secures SOTA results on 9 out of 12 tasks, surpassing even the original specialized agents (Coding, Memory, Tool).

capabilities. Traditionally, these capabilities are obtained by mixing offline datasets and performing joint training. However, training such a generalist model from scratch is computationally expensive, and maintaining separate checkpoints for each task is storage-inefficient. Consequently, model merging, which combines multiple task-specific models fine-tuned from the same base model into a unified model, has emerged as a widely adopted solution (Ilharco et al., 2023; Yadav et al., 2023; Matena and Raffel, 2022; Yu et al., 2024). It offers significant advantages, including data privacy preservation, the elimination of additional training costs, and minimal sacrificed performance.

Recently, post-training has shifted from supervised fine-tuning (SFT) to reinforcement learning (RL) (Team et al., 2025a; Guo et al., 2025), particularly for agentic models with strong reasoning

¹Code is available at [RAM](#).

capabilities. This shift fundamentally changes the role of model merging. In the SFT scenarios, multi-task capabilities can still be easily obtained by performing joint training. In contrast, joint multi-task training under on-policy RL is impractical in real-world systems, as it requires parallel task-specific environments and reward models to ensure on-policy training. Consequently, model merging becomes merely convenient solution in the RL setting. A representative example from large-scale industrial agent systems is UI-TARS2 (Wang et al., 2025a), which trains specialized vertical agents in isolated environments via RL, and subsequently merges them into a unified generalist agent. This paradigm reflects a practical compromise: specialization through RL, followed by post-hoc integration through model merging.

However, directly applying existing model merging methods to RL-trained agents leads to performance degradation. Most prior approaches, including Task Arithmetic (Ilharco et al., 2023), TIES-Merging (Yadav et al., 2023), and DARE (Yu et al., 2024), are developed under the assumption of SFT parameter updates, or task vectors, and are therefore mismatched to the RL setting. Unlike SFT, which typically induces dense and redundant parameter updates (Chu et al., 2025; Shenfeld et al., 2025), on-policy RL produces highly sparse and often disjoint task vector distributions, shaped by task-specific reward signals and RL objectives that target narrow behaviors. When such sparse updates are globally averaged during merging, task-specific unique parameter updates are divided by the number of models, resulting in signal dilution, which degrades task-specific behavior knowledge.

To address this mismatch, we propose **Reinforced Agent Merging (RAM)**, a distribution-aware merging framework designed specifically for RL-trained agents. RAM explicitly disentangles shared and task-specific unique regions of task vectors obtained via RL processes, averaging shared regions to preserve common capabilities while selectively preserving and rescaling unique regions to prevent signal dilution. By maintaining RL task vectors during merging, RAM enables specialized behavior knowledge from each model to coexist within the merged unified model. Our contributions for reinforced agent merging are:

- We identify the mismatch between merging RL-trained agentic models and existing merging method for SFT-trained models, which is

signal dilution in heterogeneous distribution of sparse parameter updates.

- We propose a distribution-aware merging method that treats shared and task-unique parameter updates differently, averaging the former while preserving and rescaling the latter based on distribution, avoiding signal dilution and compensate for performance degradation.
- Extensive experiments demonstrate that RAM not only outperforms existing merging methods across diverse architectures and domains but also unlocks synergistic potential among agents. Unified RAM model achieves performance superior to that of individual specialized agents on their domain tasks.

2 Related Works

Post-training Agents with RL RL has recently emerged as a pivotal paradigm for enhancing the reasoning capabilities of Large Language Models (LLMs) (Jaech et al., 2024; Shao et al., 2024; Team et al., 2025a). Several general-purpose algorithms, such as PPO (Schulman et al., 2017), GRPO (Guo et al., 2025), and DAPO (Yu et al., 2025b), have been developed to support this direction. Beyond general reasoning, RL is extensively applied to specialize agents for diverse domains. In coding, methods like CURE (Wang et al., 2025b), SWERL (Wei et al., 2025b), and GLM-4.5 (Zeng et al., 2025) have demonstrated significant success. For memory extension, MemAgent (Yu et al., 2025a) and various cache-based approaches (Shi et al., 2025) optimize long-context handling. Furthermore, RL has been instrumental in developing tool-integrated reasoning agents, such as AutoTIR (Wei et al., 2025a) and ToolRL (Qian et al., 2025), as well as search-augmented agents (Jin et al., 2025b; Team et al., 2025b; Sun et al., 2025) and computer-use agents (Wang et al., 2025a; Ye et al., 2025; Wanyan et al., 2025).

Model Merging for LLMs Model merging has demonstrated superiority in multi-task learning by synthesizing different task-specific models into a single entity without additional training (Ilharco et al., 2023; Jin et al., 2025a; Matena and Raffel, 2022). Techniques such as Task Arithmetic (Ilharco et al., 2023), TIES-Merging (Yadav et al., 2023), and DARE (Yu et al., 2024) have been successfully validated for merging SFT models across multiple tasks. Beyond multi-task integration, merging




Figure 2: **Left:** Sparsity of task vectors varies between agent models. **Right:** Non-zero elements distributions of task vector varies on the number of overlaps with other task vectors.

has also been employed to mitigate model collapse (IMM (Yuan et al., 2025b)) and enhance reasoning efficiency (Wu et al., 2025). While recent works like UI-TARS2 (Wang et al., 2025a) attempt to merge RL-trained agentic models, they rely on simple weight interpolation, which remains suboptimal for this regime. Distinguished from prior studies, our work is the first to systematically characterize the unique behaviors of task vectors induced by RL and to design a tailored merging strategy, that aligns their specific heterogeneity.

3 Reinforced Task Vector Behaviors

In this section, we characterize the properties of task vectors derived from reinforcement learning and explain how these properties render existing model merging methods suboptimal.

3.1 Preliminaries and Settings

Task Vectors Task vector is the set of parameter updates for a specific task. Let $\theta_{\text{pre}} \in \mathbb{R}^d$ denote the parameters of a pre-trained base model. We consider N different tasks, where each task $t \in \{1, \dots, N\}$ yields a fine-tuned model θ_t . The task vector for task t is defined as $\tau_t = \theta_t - \theta_{\text{pre}}$. In this study, we specifically examine task vectors induced by RL fine-tuning, referring to them as **Reinforced Task Vectors**. The objective of model merging is to synthesize a single merged task vector τ_{merged} to construct a final model $\theta_{\text{merged}} = \theta_{\text{pre}} + \tau_{\text{merged}}$.

Vector Sparsity We define the task vector sparsity of a model θ_t relative to the base model as $\text{sparsity}(\theta_t, \theta_{\text{pre}}) := 1 - \|\theta_t - \theta_{\text{pre}}\|_0/d$, where $\|\cdot\|_0$ represents the number of non-zero elements and d is the total parameter dimension. Following




Figure 3: The performance gain (%) of merging unique regions of reinforced task vectors across domains.

standard practice² (Mukherjee et al., 2025; Paszke et al., 2019), we consider two parameters equal (implying a zero element in the task vector) if their absolute difference is $\leq 10^{-5}$.

Reinforced Agentic Models To investigate these properties, we select three representative agentic models trained via RL to specialize in coding, tool-use, and long-context memory as follows.

CURE (Wang et al., 2025b): A coding agent that co-evolves with a unit tester to enhance code generation capabilities. **ToolRL** (Qian et al., 2025): A reasoning agent optimized for general-purpose tool selection and application. **MemAgent** (Yu et al., 2025a): An agent optimized for long-context tasks, with a workflow for extended memory retention. All reinforced agents are initialized from the same base model, Qwen2.5-7B-Instruct (Yang et al., 2024).

To evaluate their specialized agentic capabilities and reinforced task vector behaviors, we employ benchmarks across the coding, tool-use, and memory domains; detailed evaluation settings are provided in Section 5.1.

3.2 Heterogeneity in Reinforced Task Vectors

In this section, we have the two key observations by analyzing the distribution of reinforced task vectors in parameter space: 1) Reinforced task vectors are sparsely distributed in parameter space, with notable heterogeneous sparsity and distributions. 2) Heterogeneity causes shared and unique regions of reinforced task vectors. Unique regions are critical for improving corresponding agentic domain performances, and exert little negative interference on other domains' performances.

²PyTorch uses 10^{-5} as the default tolerance for gradient checking (refer to PyTorch Documentation).

226 **Heterogeneity in Sparsity and Distribution** Recent work (Mukherjee et al., 2025; Yuan et al.,
227 2025a) highlights the inherent sparsity of RL up-
228 dates: unlike SFT, which updates global pa-
229 rameters, RL tends to fine-tune specific sub-networks.
230 Our analysis confirms this and further unveils het-
231 erogeneous patterns in both sparsity and spatial
232 distribution of them.

233 First, the percentage of non-zero elements, or the
234 sparsity levels of reinforced task vectors vary dras-
235 tically. As shown in Figure 2 (left), the coding agent
236 exhibits extreme sparsity, modifying only 3.2% of
237 parameters. In contrast, agents optimized for tool
238 use and long-context memory induce significantly
239 denser updates, affecting 46.2% and 54.3% of the
240 parameter space, respectively.

241 Second, these non-zero elements are distributed
242 across disparate regions, creating different spatial
243 overlap patterns. We categorize parameters (task
244 vector elements) based on whether they are up-
245 dated by a single agent (unique) or multiple agents
246 (shared). Figure 2 (right) reveals that the coding,
247 tool-use, and memory agents concentrate differ-
248 ent fractions of their updates in unique, non-
249 overlapping regions (6.3%, 40.8%, and 47.5% of
250 their respective non-zero elements). Analysis in
251 Appendix C.1 confirms that this heterogeneity gen-
252 eralizes to other reinforced agents with different
253 model architectures and specialized domains.

255 **The Role of Heterogeneity** We further investi-
256 giate the impact of these heterogeneous, unique
257 regions on both in-domain and out-of-domain agen-
258 tic tasks. In our experiments, we isolate the unique
259 component of a single task vector, merge it into
260 the base model, and evaluate performance across
261 all domains. Results in Figure 3 demonstrate that
262 unique regions cause almost no negative interfer-
263 ence on out-of-domain tasks (occasionally yielding
264 improvements) while driving significant gains on
265 the corresponding in-domain tasks. When we in-
266 tentionally dilute the magnitude of these unique
267 vectors to $1/N$ ($N = 3$) of their original value, in-
268 domain performance drops sharply, verifying the
269 positive contribution of these unique components.

270 **Signal Dilution** Previous heterogeneity analysis
271 uncovers the root cause of performance degrada-
272 tion when using previous SOTA merging meth-
273 ods on RL models. Existing methods typically
274 employ element-wise averaging or variants (e.g.,
275 $\tau_{\text{merged}} = \frac{1}{N} \sum \tau_i$). While averaging is benefi-
276 cial for shared regions of task vectors by balancing

277 multi-task performance, it is detrimental to task-
278 specific, unique regions in RL scenarios. For a
279 unique parameter updates for task t , averaging
280 with $N - 1$ zero-valued updates effectively scales
281 its magnitude by $1/N$. This operation dilutes the
282 learned signal without providing any balancing ben-
283 efit, as these regions hardly interfere with out-of-
284 domain tasks. We term this phenomenon *Signal*
285 *Dilution*. Figure 3 illustrates that this dilution (sim-
286 ulated with $N = 3$) causes significant performance
287 regression. The prevalence of unique task vector
288 components in RL agents therefore necessitates a
289 merging strategy capable of disentangling these
290 regions to prevent signal dilution. We provide a
291 detailed analysis of signal dilution in each existing
292 merging strategy in Appendix D.

4 Merging Reinforced Agentic Models

293 Our analysis in Section 3.2 demonstrates that rein-
294 forced task vectors are inherently sparse and het-
295 erogeneous. While task vectors or parameters up-
296 dated by multiple agents (shared regions) benefit
297 from averaging to stabilize the consensus direc-
298 tion, parameter updated by a single agent (unique
299 regions) suffer from *Signal Dilution* when stan-
300 dard averaging is applied. To address this, we
301 propose **Reinforced Agent Merging (RAM)** illus-
302 trated in Figure 4, a method that explicitly disen-
303 tangles these regions based on distribution statis-
304 tics of task vectors. RAM applies selective merg-
305 ing strategies: it averages shared parameters to
306 absorb unified multi-task capabilities while pre-
307 serving the full magnitude of unique parameters
308 to prevent signal dilution. Additionally, we intro-
309 duce a distribution-aware rescaling mechanism to
310 further amplify unique task capabilities.

4.1 Probing Vector Distribution

312 First, we probe the active updated parameters for
313 each reinforced task vector τ_t . Using the threshold
314 established in Section 3.2 (e.g., $\epsilon = 10^{-5}$), we
315 compute a binary mask $\mathbf{m}_t \in \{0, 1\}^d$ for task t :

$$m_{t,i} = \mathbb{I}(|\tau_{t,i}| > \epsilon), \quad (1)$$

318 where i indexes the parameter dimensions and
319 $\mathbb{I}(\cdot)$ is the indicator function. We then define
320 the overlap count vector $\mathbf{c} = \sum_{t=1}^N \mathbf{m}_t$, where
321 $c_i \in \{0, \dots, N\}$ represents the number of agents
322 that actively update the i -th parameter.

323 For any given task t , the set of total updated
324 parameters is partitioned into two disjoint subsets:




Figure 4: **Method Overview.** (a) A base model is trained via RL to different agents, we track the distributions of obtained reinforced task vectors. (b) Probing the distribution of task vectors to shared, unique, unchanged sets. (c) selective merging task vectors by averaging shared set and rescaling unique set to the base model.

the **Shared** subset (where $c_i \geq 2$) and the **Unique** subset (where $c_i = 1$). To quantify the structural distribution of the agent’s updates, we define the *Overlap-Unique Ratio* ρ_t :

$$\rho_t = \frac{\sum_{i:c_i \geq 2} m_{t,i}}{\sum_{i:c_i=1} m_{t,i}}. \quad (2)$$

Here, the numerator represents the count of shared parameters, and the denominator represents the count of unique parameters. Since the sum of these two components constitutes the total updated parameter volume, a higher ρ_t indicates that model learns task t largely within the shared subspace.

4.2 Rescaling Unique Regions

Reinforced task vectors with high Overlap-Unique Ratios (ρ_t) are likely to suffer greater degradation of task capabilities when their substantial shared regions are averaged with other vectors. Therefore, we proportionally rescale the unique regions of such parts of task vectors to compensate for the performance loss incurred in the shared regions. To achieve this, we calculate a task-specific scaling factor λ_t derived from a functional equivalence hypothesis as follows.

Let Δf_t denote the functional gain, or task performance gain, induced by the task vector τ_t . We decompose the total gain into shared and unique components based on the overlap statistics defined in Section 4.1:

$$\Delta f_t = \mathcal{C}_{\text{shared}} + \mathcal{C}_{\text{unique}}, \quad (3)$$

where the gains are modeled as the task vector element $\tau_{t,k}$ weighted by local sensitivity g_i , which indicates the contribution coefficient

mapping the task vector element to performance gain. The performance gains are therefore represented as: $\mathcal{C}_{\text{shared}} = \sum_{i:c_i \geq 2} g_i \tau_{t,i} m_{t,i}$, $\mathcal{C}_{\text{unique}} = \sum_{i:c_i=1} g_i \tau_{t,i} m_{t,i}$. In the merged task vectors, elements in the shared regions are averaged across vectors, and the corresponding task performances are degraded. We model this as a contraction of the effective signal by a coefficient $1 - r \in (0, 1)$. To counteract this, we rescale the magnitudes of parameters in unique regions and have a new functional gain expression:

$$\hat{\Delta f}_t = (1 - r)\mathcal{C}_{\text{shared}} + \lambda_t \mathcal{C}_{\text{unique}}. \quad (4)$$

We hypothesize that this rescaling operation achieves the *functional equivalence* to have the same performance gain for the task vector τ_t on task t : $\hat{\Delta f}_t \approx \Delta f_t$. Under the simplifying assumption that parameter importance is isotropic ($g_i \tau_{t,i} \approx \text{const}$ on average), the ratio of functional contributions $\mathcal{C}_{\text{shared}}/\mathcal{C}_{\text{unique}}$ reduces to the ratio of parameter counts defined in Eq. 2. Then by solving *functional equivalence* as our objective, the required amplification satisfies:

$$\lambda_t - 1 = r \frac{\mathcal{C}_{\text{shared}}}{\mathcal{C}_{\text{unique}}} \approx r \frac{\sum_{i:c_i \geq 2} m_{t,i}}{\sum_{i:c_i=1} m_{t,i}} = r \rho_t, \quad (5)$$

which is approximately proportional to the Overlap-Unique Ratio ρ_t . This suggests that tasks with higher overlap require stronger compensation in their unique parameter subspace to counteract the degradation induced by averaging. However, directly instantiating this proportional relationship may lead to numerical instability when ρ_t is large. To balance signal compensation with stability, we employ a clipped linear scaling rule:

$$\lambda_t = 1 + r \cdot \text{clip}(\rho_t, 0, \alpha), \quad (6)$$

389 where r controls the overall amplification strength
 390 and α serves as a stable bound. This design pre-
 391 serves the monotonic growth implied by the hy-
 392 pothesis while preventing excessive amplification
 393 in high-overlap scenarios.

394 4.3 Selective Merging

395 Finally, we construct the merged task vector
 396 τ_{merged} element-wise. Unlike existing merging
 397 methods, which effectively divide unique par-
 398 ameters by N (causing signal dilution), our strategy
 399 differentiates between shared and unique regions.
 400 For each parameter index i , let $\mathcal{T}_i = \{t \mid m_{t,i} = 1\}$
 401 denote the set of indices of active tasks for that pa-
 402 rameter. Note that the cardinality $|\mathcal{T}_i|$ corresponds
 403 to the overlap count c_i defined in Section 4.1. The
 404 merged element $\tau_{\text{merged},i}$ is computed as:

$$\tau_{\text{merged},i} = \begin{cases} 0 & \text{if } |\mathcal{T}_i| = 0, \\ \lambda_t \cdot \tau_{t,i} & \text{if } \mathcal{T}_i = \{t\}, \\ \frac{1}{|\mathcal{T}_i|} \sum_{t \in \mathcal{T}_i} \tau_{t,i} & \text{if } |\mathcal{T}_i| \geq 2. \end{cases} \quad (7)$$

405 This selective strategy ensures that: 1) *Shared
 406 Knowledge* ($|\mathcal{T}_i| \geq 2$) is averaged to balance
 407 multi-task capabilities. 2) *Task-Specific Knowledge*
 408 ($|\mathcal{T}_i| = 1$) is completely preserved and amplified by
 409 λ_t to compensate for the contraction of the effective
 410 signal in shared regions, explicitly targeting func-
 411 tional equivalence. 3) *No Knowledge* ($|\mathcal{T}_i| = 0$) is
 412 set to zero to filter out insignificant parameter fluc-
 413 tuations and ensure that the base model’s general
 414 capabilities remain undisturbed.

416 5 Experiments

417 5.1 Setup

418 **Baselines** We categorize our baselines into two
 419 groups: original specialized agent models and es-
 420 tablished model merging techniques. **(i) Original
 421 Agent Models:** We utilize **Qwen2.5-7B**-
 422 **Instruct** (Yang et al., 2024) as the shared base
 423 model. The task-specific reinforced agents in-
 424 clude: **CURE** (Wang et al., 2025b) for coding,
 425 **ToolRL** (Qian et al., 2025) for tool-use, and
 426 **MemAgent** (Yu et al., 2025a) for long-context
 427 memory. **(ii) Merging Methods:** We compare
 428 our approach against prominent merging strate-
 429 gies: **Task Arithmetic** (Ilharco et al., 2023), which
 430 linearly combines task vectors; **Fisher Merg-
 431 ing** (Matena and Raffel, 2022), which weighs
 432 parameters based on Fisher information; **TIES-
 433 Merging** (Yadav et al., 2023), which mitigates
 434 parameter interference through trimming and sign

435 consensus; and **DARE** (Yu et al., 2024), which
 436 randomly drops and rescales the parameters. Fol-
 437 lowing standard practice, we combine DARE with
 438 Task Arithmetic and TIES for evaluation. We intro-
 439 duce more details about baselines in Appendix D.

440 **Evaluations** We evaluate the models across three
 441 critical agentic domains: coding, tool-use, and
 442 long-context memory. For **Coding**, we measure
 443 generated code pass accuracy (ACC) and unit test
 444 pass accuracy (UT) on the LiveBench (White et al.,
 445 2025) and LiveCodeBench (Jain et al., 2024) bench-
 446 marks. For **Tool Use**, we utilize the Berkeley Func-
 447 tion Call Leaderboard (BFCL) (Patil et al., 2025),
 448 specifically reporting results on the Live/Non-Live
 449 Parallel (Para) and Parallel Multiple (P_Mul) sub-
 450 sets. For **Long-Context Memory**, following (Yu
 451 et al., 2025a), we employ the RULER bench-
 452 mark (Hsieh et al., 2024) to assess performance
 453 on long-context tasks, including RULER (Hsieh
 454 et al., 2024) HotpotQA and SQuAD. Additional
 455 datasets and detailed evaluation criteria are pro-
 456 vided in Appendix B.

457 **Implementations** For hyperparameters, we set
 458 $r = 0.1$ and $\alpha = 2.0$. We denote our method
 459 without task-specific rescaling as **RAM** and with
 460 rescaling as **RAM+**. RAM is the special case of
 461 RAM+ when $r = 0$. We provide the details of
 462 pseudocode and agents for merging in Appendix A.

463 5.2 RAM is a Better Fit for Reinforced Models

464 Table 1 demonstrates that both RAM and RAM+
 465 consistently outperform all baselines, establishing
 466 a new SOTA. Specifically, RAM achieves an aver-
 467 age score of 64.82 across all tasks, surpassing the
 468 strongest baseline DARE (63.33). Building on this
 469 foundation, RAM+ further pushes the boundary to
 470 66.55 after rescaling unique regions, inducing extra
 471 improvements where the merged generalist exceeds
 472 the capabilities of specialized task agents on most
 473 of evaluations. For instance, in the **Coding** domain,
 474 RAM+ surpasses the specialist Coding agent on
 475 LiveBench and LiveCodeBench, suggesting that
 476 reasoning signals from other tasks enhance coding
 477 precision. This superiority extends to **Tool Use**,
 478 where RAM+ significantly outperforms the Tool
 479 agent in complex parallel scenarios (Live P_Mul:
 480 70.83 vs. 58.33), and to **Long-Context Memory**,
 481 where it achieves global optimal performance on
 482 SQuAD 64k (82.03), beating the dedicated Mem-
 483 ory agent (77.34).

Model	Coding				Tool Using				Memory				Avg
	LiveBench		LiveCodeBench		Live		Non-Live		HotpotQA		SQuAD		
	ACC	UT	ACC	UT	Para	P_Mul	Para	P_Mul	7K	14K	32K	64K	
<i>Base and Task Models</i>													
Base	28.35	40.87	23.43	36.42	56.25	41.67	68.00	55.00	60.94	50.00	64.84	58.59	48.70
CURE (Coding)	37.70	49.27	30.23	45.76	56.25	37.50	64.00	51.50	58.59	56.25	60.94	44.22	49.35
ToolRL (Tool)	31.84	41.36	26.76	42.05	56.25	58.33	91.00	89.00	58.59	48.44	59.38	46.95	54.16
MemAgent (Memory)	39.25	50.12	28.92	44.80	37.50	50.00	78.50	48.50	78.91	78.12	81.25	77.34	57.77
<i>Merged Models</i>													
TA	38.09	<u>51.62</u>	<u>31.95</u>	46.69	43.75	45.83	87.50	69.50	69.53	68.75	73.44	72.66	58.28
Fisher	36.72	48.73	30.87	45.89	43.75	41.67	86.5	63.5	60.06	49.22	58.59	60.94	52.20
TIES	<u>39.25</u>	49.88	30.63	46.32	43.75	54.17	84.00	67.50	71.88	82.03	75.00	75.78	60.02
DARE+TA	37.50	48.60	<u>31.95</u>	46.69	50.00	62.50	92.50	89.50	76.56	76.56	<u>77.34</u>	70.31	63.33
DARE+TIES	35.93	45.66	29.26	39.53	56.25	58.33	91.50	90.00	75.00	77.34	76.56	74.22	62.47
RAM	38.28	49.71	31.96	47.72	56.25	66.67	91.00	91.50	75.78	75.00	74.22	79.69	64.82
RAM+	40.23	52.57	31.60	46.84	56.25	70.83	90.50	91.00	79.69	78.13	78.91	82.03	66.55

Table 1: **Main results of agent merging.** We evaluate the capabilities across three domains: Coding (LiveBench, LiveCodeBench), Tool Use (Live, Non-Live), and Memory (HotpotQA, SQuAD). **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across *all evaluated models*, including the specialized Task Models.




Figure 5: The performances of merging two agents across domains.

5.3 Extending Model Combinations

To evaluate the effectiveness of RAM beyond tri-agent merging, we extend pairwise agent merging experiments with more model combinations across Tool+Memory, Coding+Tool, and Coding+Memory scenarios, as illustrated in Figure 5 (details are provided in Appendix C). Across all three combinations, RAM/RAM+ consistently achieves the highest average performance, demonstrating superior robustness on various combinations compared to baselines. Specifically, in the Coding+Tool setting, RAM+ attains an average score of 60.04, significantly outperforming the strongest baseline DARE+TIES (56.74) and effectively bridging the capability gap that tradi-

tional methods like Task Arithmetic and TIES fail to address due to signal dilution. Similarly, in the Tool+Memory and Coding+Memory scenarios, RAM+ maintains dominant performance with average scores of 75.86 and 61.21 respectively, confirming that RAM can be successful in multiple agent combinations.

5.4 Ablation Study

We ablate our method to RAM ($r = 0$) and RAM+, and investigate the sensitivity of our proposed method to the scaling factor r . Table 2 presents the ablation results with r varying from 0.00 to 0.20. Note that when $r = 0$, the method represents RAM. As observed, the model's overall per-

484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505

499
500
501
502
503
504
505

r	Code				Tool				Memory				Avg	
	LiveBench		LiveCodeBench		Live		Non-Live		HotpotQA		RulerQA			
	ACC	UT	ACC	UT	Para	P_Mul	Para	P_Mul	7k	14k	32k	64k		
0.00	38.28	49.71	<u>31.96</u>	47.72	<u>56.25</u>	<u>66.67</u>	91.00	91.50	75.78	75.00	74.22	<u>79.69</u>	64.82	
0.05	37.70	<u>50.49</u>	30.63	45.59	62.50	62.50	92.00	91.50	75.78	79.69	75.00	82.03	65.45	
0.10	40.23	52.57	31.60	46.84	<u>56.25</u>	70.83	90.50	<u>91.00</u>	79.69	<u>78.13</u>	78.91	82.03	66.55	
0.15	38.67	49.85	31.41	46.53	62.50	66.67	89.50	91.50	<u>78.12</u>	79.69	<u>79.69</u>	75.78	<u>65.83</u>	
0.20	39.45	50.36	32.58	47.54	<u>56.25</u>	62.50	91.00	90.50	<u>78.12</u>	78.12	80.47	77.34	65.35	

Table 2: **Ablation Study.** **Bold** and underlined values denote the best and second-best performance.




Figure 6: Merging results for RL agents trained from Llama3.2-3B-Instruction base model.

513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580

formance (Avg) exhibits a trend of initially increasing and then decreasing. The performance peaks at $r = 0.10$, achieving the highest average score of 66.55. Specifically, setting $r = 0.10$ (RAM+) yields the best or second-best results across the majority of metrics, particularly showing significant gains in LiveBench (Coding) and HotpotQA (Memory) compared to RAM ($r = 0.00$). However, further increasing r beyond 0.10 leads to diminishing returns, with the average score dropping to 65.35 at $r = 0.20$. This suggests that while a moderate scaling factor effectively enhances task-specific capabilities, an excessively large r may disrupt the general knowledge of the merged model.

5.5 Extending Architecture and Domains

Besides the agents trained from Qwen architecture, we extend the experiment to Llama architecture and additional domains. We choose Llama3.2-3B (Grattafiori et al., 2024) as the base model, and choose models trained from it via RL: search agent ZeroSearch (Sun et al., 2025), math reasoning agent (Zhao et al., 2025), and tool-using agent ToolRL (Qian et al., 2025). The evaluation details are provided in Appendix B.4. Figure 6 illustrates that, consistent with our findings on the Qwen, RAM and RAM+ demonstrate superior performance across multiple agentic domains, consistently outperforming baselines. Notably, they achieve positive synergy in both Math and Tool

domains, where the merged generalist surpasses the performance of the original specialized agents. For instance, in the Tool domain, RAM+ exhibits a significant margin over the specialist, suggesting that reasoning capabilities from Math and Search agents synergize to enhance tool-use. In the Search domain, RAM/RAM+ successfully retain original capability, whereas baselines show notable regression. These results confirm that the heterogeneity of reinforced task vectors is a general property, and RAM effectively addresses this by preserving task-specific specialized knowledge independent of model scale and architecture.

Additional Experiments Besides the above experiments, we further provide the instruction following evaluation to assess the forgetting after merging in Appendix C.2; Merging efficiency comparison in Appendix C.3; Evaluation on the additional tasks in Appendix C.5; Additional rescaling strategy performance in Appendix E.

6 Conclusion

In this work, we address the critical challenge of merging agents fine-tuned via RL, identifying a fundamental mismatch between standard merging techniques designed for dense SFT updates and the sparse, heterogeneous nature of on-policy RL task vectors. We demonstrate that treating global task vectors equally in previous methods in this setting leads to signal dilution of task-specific capabilities. To bridge this gap, we propose Reinforced Agent Merging (RAM), a framework that explicitly disentangles shared and unique parameter update regions and applies a distribution-based rescaling strategy to preserve specialized behaviors. Extensive evaluations across multiple agentic domains and model architectures show that RAM significantly outperforms existing baselines, achieving SOTA results and surpassing the original specialists as a unified generalist on most tasks.

581 Limitations

582 While Reinforced Agent Merging (RAM) effec-
583 tively mitigates signal dilution for RL-trained
584 agents, our current study has certain boundaries.
585 First, our experiments focus on merging a com-
586 mon number of agents; as N scales significantly,
587 the probability of parameter collision in the shared
588 subspace increases, potentially requiring more com-
589 plex conflict resolution strategies beyond simple
590 averaging. Second, the derivation of our rescaling
591 factor relies on an isotropic assumption of param-
592 eter importance, which, while empirically robust,
593 does not explicitly account for element-wise cur-
594 vature information that could offer finer-grained
595 control at a higher computational cost. Third,
596 although we identified a default hyperparameter
597 configuration that generalizes well across Qwen
598 and Llama architectures, optimal performance on
599 agents trained with fundamentally different RL al-
600 gorithms or modalities may require task-specific
601 tuning. Finally, our evaluation is primarily con-
602 ducted on 3B and 7B parameter models; verifying
603 whether the sparsity hypothesis and RAM’s
604 efficacy persist in massive-scale models (70B+)
605 remains an open question for future research.

606 References

607 Jacob Austin, Augustus Odena, Maxwell Nye, Maarten
608 Bosma, Henryk Michalewski, David Dohan, Ellen
609 Jiang, Carrie Cai, Michael Terry, Quoc Le, and 1
610 others. 2021. Program synthesis with large language
611 models. *arXiv preprint arXiv:2108.07732*.

612 Tianzhe Chu, Yuexiang Zhai, Jihan Yang, Shengbang
613 Tong, Saining Xie, Dale Schuurmans, Quoc V Le,
614 Sergey Levine, and Yi Ma. 2025. SFT memorizes,
615 RL generalizes: A comparative study of foundation
616 model post-training. In *Forty-second International
617 Conference on Machine Learning*.

618 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian,
619 Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias
620 Plappert, Jerry Tworek, Jacob Hilton, Reiichiro
621 Nakano, and 1 others. 2021. Training verifiers
622 to solve math word problems. *arXiv preprint
623 arXiv:2110.14168*.

624 Leo Gao, Jonathan Tow, Baber Abbasi, Stella Bider-
625 man, Sid Black, Anthony DiPofi, Charles Foster,
626 Laurence Golding, Jeffrey Hsu, Alain Le Noac’h,
627 Haonan Li, Kyle McDonell, Niklas Muennighoff,
628 Chris Ociepa, Jason Phang, Laria Reynolds, Hailey
629 Schoelkopf, Aviya Skowron, Lintang Sutawika, and
630 5 others. 2024. [The language model evaluation harness](#).

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al- Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, and 1 others. 2024. The llama 3 herd of models. <i>arXiv preprint arXiv:2407.21783</i> .	632
Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shi- rong Ma, Peiyi Wang, Xiao Bi, and 1 others. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. <i>arXiv preprint arXiv:2501.12948</i> .	637
Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the math dataset. In <i>Thirty- fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track</i> .	643
Xanh Ho, Anh-Khoa Duong Nguyen, Saku Sugawara, and Akiko Aizawa. 2020. Constructing a multi-hop qa dataset for comprehensive evaluation of reasoning steps. <i>arXiv preprint arXiv:2011.01060</i> .	649
Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shan- tanu Acharya, Dima Rekesh, Fei Jia, Yang Zhang, and Boris Ginsburg. 2024. Ruler: What’s the real context size of your long-context language models? <i>arXiv preprint arXiv:2404.06654</i> .	653
Gabriel Ilharco, Marco Túlio Ribeiro, Mitchell Worts- man, Ludwig Schmidt, Hannaneh Hajishirzi, and Ali Farhadi. 2023. Editing models with task arithmetic. In <i>The Eleventh International Conference on Learn- ing Representations</i> .	658
Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richard- son, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, Alex Carney, and 1 others. 2024. Openai o1 system card. <i>arXiv preprint arXiv:2412.16720</i> .	663
Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar- Lezama, Koushik Sen, and Ion Stoica. 2024. Live- codebench: Holistic and contamination free eval- uation of large language models for code. <i>arXiv preprint arXiv:2403.07974</i> .	668
Bowen Jin, Hansi Zeng, Zhenrui Yue, Jinsung Yoon, Sercan Arik, Dong Wang, Hamed Zamani, and Jiawei Han. 2025a. Search-r1: Training llms to reason and leverage search engines with reinforcement learning. <i>arXiv preprint arXiv:2503.09516</i> .	674
Bowen Jin, Hansi Zeng, Zhenrui Yue, Jinsung Yoon, Sercan O Arik, Dong Wang, Hamed Zamani, and Jiawei Han. 2025b. Search-r1: Training LLMs to reason and leverage search engines with reinforce- ment learning. In <i>Second Conference on Language Modeling</i> .	679
Tom Kwiatkowski, Jennimaria Palomaki, Olivia Red- field, Michael Collins, Ankur Parikh, Chris Alberti,	685
	686

687	Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, and 1 others. 2019. Natural questions: a benchmark for question answering research. <i>Transactions of the Association for Computational Linguistics</i> .	741
688		742
689		743
690		744
691		
692		745
693	Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with pagedattention. In <i>Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles</i> .	746
694		747
695		748
696		749
697		750
698		
699		751
700	Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, and 1 others. 2022. Competition-level code generation with alphacode. <i>Science</i> .	752
701		753
702		754
703		755
704		
705	Michael S Matena and Colin A Raffel. 2022. Merging models with fisher-weighted averaging. <i>Advances in Neural Information Processing Systems</i> .	
706		
707		756
708	Sagnik Mukherjee, Lifan Yuan, Dilek Hakkani-Tur, and Hao Peng. 2025. Reinforcement learning finetunes small subnetworks in large language models. <i>arXiv preprint arXiv:2505.11711</i> .	757
709		758
710		759
711		760
712	Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, and 1 others. 2022. Training language models to follow instructions with human feedback. <i>Advances in neural information processing systems</i> .	761
713		
714		
715		
716		
717		762
718	Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, and 1 others. 2019. Pytorch: An imperative style, high-performance deep learning library. <i>Advances in neural information processing systems</i> , 32.	763
719		764
720		765
721		766
722		
723		
724		767
725	Shishir G Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. 2025. The berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In <i>Forty-second International Conference on Machine Learning</i> .	768
726		769
727		770
728		771
729		
730		772
731	Guilherme Penedo, Anton Lozhkov, Hynek Kydlíček, Loubna Ben Allal, Edward Beeching, Agustín Piquerres Lajarín, Quentin Gallouédec, Nathan Habib, Lewis Tunstall, and Leandro von Werra. 2025. Codeforces. https://huggingface.co/datasets/open-r1/codeforces .	773
732		774
733		775
734		776
735		
736		777
737	Cheng Qian, Emre Can Acikgoz, Qi He, Hongru WANG, Xiusi Chen, Dilek Hakkani-Tür, Gokhan Tur, and Heng Ji. 2025. ToolRL: Reward is all tool learning needs. In <i>The Thirty-ninth Annual Conference on Neural Information Processing Systems</i> .	778
738		779
739		780
740		781
741		782
742	John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. <i>arXiv preprint arXiv:1707.06347</i> .	783
743		784
744		785
745	Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, and 1 others. 2024. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. <i>arXiv preprint arXiv:2402.03300</i> .	786
746		787
747		788
748		789
749		790
750		791
751		792
752	Idan Shenfeld, Jyothish Pari, and Pulkit Agrawal. 2025. RL's razor: Why on-policy reinforcement learning forgets less. In <i>AI That Keeps Up: NeurIPS 2025 Workshop on Continual and Compatible Foundation Model Updates</i> .	793
753		
754		
755		
756	Dachuan Shi, Yonggan Fu, Xiangchi Yuan, Zhongzhi Yu, Haoran You, Sixu Li, Xin Dong, Jan Kautz, Pavlo Molchanov, and Yingyan Celine Lin. 2025. Lacache: Ladder-shaped kv caching for efficient long-context modeling of large language models. In <i>Forty-second International Conference on Machine Learning</i> .	794
757		795
758		796
759		797
760		798
761		
762	Hao Sun, Zile Qiao, Jiayan Guo, Xuanbo Fan, Yingyan Hou, Yong Jiang, Pengjun Xie, Fei Huang, and Yan Zhang. 2025. Zerosearch: Incentivize the search capability of llms without searching. <i>arXiv preprint arXiv:2505.04588</i> .	799
763		800
764		801
765		802
766		
767	Kimi Team, Angang Du, Bofei Gao, Bowei Xing, Changjiu Jiang, Cheng Chen, Cheng Li, Chenjun Xiao, Chenzhuang Du, Chonghua Liao, and 1 others. 2025a. Kimi k1. 5: Scaling reinforcement learning with llms. <i>arXiv preprint arXiv:2501.12599</i> .	803
768		804
769		805
770		806
771		
772	Tongyi DeepResearch Team, Baixuan Li, Bo Zhang, Dingchu Zhang, Fei Huang, Guangyu Li, Guoxin Chen, Huifeng Yin, Jialong Wu, Jingren Zhou, and 1 others. 2025b. Tongyi deepresearch technical report. <i>arXiv preprint arXiv:2510.24701</i> .	807
773		808
774		809
775		810
776		
777	Haoming Wang, Haoyang Zou, Huatong Song, Jiazhan Feng, Junjie Fang, Junting Lu, Longxiang Liu, Qinyu Luo, Shihao Liang, Shijue Huang, and 1 others. 2025a. Ui-tars-2 technical report: Advancing gui agent with multi-turn reinforcement learning. <i>arXiv preprint arXiv:2509.02544</i> .	811
778		812
779		813
780		814
781		815
782		
783	Yinjie Wang, Ling Yang, Ye Tian, Ke Shen, and Mengdi Wang. 2025b. CURE: Co-evolving coders and unit testers via reinforcement learning. In <i>The Thirty-ninth Annual Conference on Neural Information Processing Systems</i> .	816
784		817
785		818
786		819
787		
788	Yuyang Wanyan, Xi Zhang, Haiyang Xu, Haowei Liu, Junyang Wang, Jiabo Ye, Yutong Kou, Ming Yan, Fei Huang, Xiaoshan Yang, and 1 others. 2025. Look before you leap: A gui-critic-r1 model for pre-operative error diagnosis in gui automation. <i>arXiv preprint arXiv:2506.04614</i> .	820
789		821
790		822
791		823
792		824
793		

794	Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2021. Finetuned language models are zero-shot learners. <i>arXiv preprint arXiv:2109.01652</i> .	849
795		850
796		851
797		852
798		853
799	Yifan Wei, Xiaoyan Yu, Yixuan Weng, Tengfei Pan, Angsheng Li, and Li Du. 2025a. Autotir: Autonomous tools integrated reasoning via reinforcement learning. <i>arXiv preprint arXiv:2507.21836</i> .	854
800		855
801		856
802		857
803	Yuxiang Wei, Olivier Duchenne, Jade Copet, Quentin Carboneaux, Lingming Zhang, Daniel Fried, Gabriel Synnaeve, Rishabh Singh, and Sida I Wang. 2025b. Swe-rl: Advancing llm reasoning via reinforcement learning on open software evolution. <i>arXiv preprint arXiv:2502.18449</i> .	858
804		859
805		860
806		861
807		862
808		863
809	Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Benjamin Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Sreemanti Dey, Shubh-Agrawal, Sandeep Singh Sandha, Siddartha Venkat Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. 2025. Livebench: A challenging, contamination-limited LLM benchmark. In <i>The Thirteenth International Conference on Learning Representations</i> .	864
810		865
811		866
812		867
813		868
814		869
815		870
816		871
817		872
818	Taiqiang Wu, Runming Yang, Tao Liu, Jiahao Wang, and Ngai Wong. 2025. Revisiting model interpolation for efficient reasoning. <i>arXiv preprint arXiv:2510.10977</i> .	873
819		874
820		875
821		876
822	Prateek Yadav, Derek Tam, Leshem Choshen, Colin Raffel, and Mohit Bansal. 2023. Resolving interference when merging models. <i>arXiv preprint arXiv:2306.01708</i> , 1.	877
823		878
824		879
825		880
826	An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, Guantong Dong, Haoran Wei, Huan Lin, Jialong Tang, Jialin Wang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Ma, and 40 others. 2024. Qwen2 technical report. <i>arXiv preprint arXiv:2407.10671</i> .	881
827		882
828		883
829		884
830		885
831		886
832		887
833	Jiabo Ye, Xi Zhang, Haiyang Xu, Haowei Liu, Junyang Wang, Zhaoqing Zhu, Ziwei Zheng, Feiyu Gao, Junjie Cao, Zhengxi Lu, and 1 others. 2025. Mobile-agent-v3: Fundamental agents for gui automation. <i>arXiv preprint arXiv:2508.15144</i> .	888
834		889
835		890
836		891
837		892
838	Hongli Yu, Tinghong Chen, Jiangtao Feng, Jiangjie Chen, Weinan Dai, Qiying Yu, Ya-Qin Zhang, Wei-Ying Ma, Jingjing Liu, Mingxuan Wang, and 1 others. 2025a. Memagent: Reshaping long-context llm with multi-conv rl-based memory agent. <i>arXiv preprint arXiv:2507.02259</i> .	893
839		894
840		895
841		896
842		897
843		898
844	Le Yu, Bowen Yu, Haiyang Yu, Fei Huang, and Yongbin Li. 2024. Language models are super mario: Absorbing abilities from homologous models as a free lunch. In <i>Forty-first International Conference on Machine Learning</i> .	899
845		900
846		901
847		902
848		903

883 A Implementation Details

884 A.1 Pseudocode

885 Here we provide the Pseudocode of RAM in Algorithm 1.

886 Algorithm 1 Reinforced Agent Merging (RAM)

```

Require: Task vectors  $\{\tau_t\}_{t=1}^N$ ; threshold  $\epsilon$ ; rescale strength
          $r$ ; clip bound  $\alpha$ 
Ensure: Merged task vector  $\tau_{\text{merged}}$ 
1: // Stage 1: Probing Vector Distribution (Sec. 4.1)
2: for  $t = 1$  to  $N$  do
3:    $\mathbf{m}_t \leftarrow \mathbb{I}(|\tau_t| > \epsilon)$  // Compute binary mask (Eq. 1)
4: end for
5:  $\mathbf{c} \leftarrow \sum_{t=1}^N \mathbf{m}_t$  // Compute overlap count vector
6: // Stage 2: Rescaling Unique Regions (Sec. 4.2)
7: for  $t = 1$  to  $N$  do
8:    $N_{\text{shared}} \leftarrow \sum_{i:c_{i,t} \geq 2} m_{t,i}$ 
9:    $N_{\text{unique}} \leftarrow \sum_{i:c_{i,t}=1} m_{t,i}$ 
10:   $\rho_t \leftarrow N_{\text{shared}} / N_{\text{unique}}$  // Ratio (Eq. 2)
11:   $\lambda_t \leftarrow 1 + r \cdot \text{clip}(\rho_t, 0, \alpha)$  // Task-specific scaling
12: end for
13: // Stage 3: Selective Merging (Sec. 4.3)
14: for each parameter index  $i$  do
15:    $\mathcal{T}_i \leftarrow \{t \mid m_{t,i} = 1\}$  // Set of active tasks
16:   if  $|\mathcal{T}_i| = 0$  then
17:      $\tau_{\text{merged},i} \leftarrow 0$ 
18:   else if  $|\mathcal{T}_i| = 1$  then
19:     Let  $t$  be the unique element in  $\mathcal{T}_i$ 
20:      $\tau_{\text{merged},i} \leftarrow \lambda_t \cdot \tau_{t,i}$  // Preserve & amplify unique
21:   else
22:      $\tau_{\text{merged},i} \leftarrow \frac{1}{|\mathcal{T}_i|} \sum_{t \in \mathcal{T}_i} \tau_{t,i}$  // Average shared
23:   end if
24: end for
25: return  $\tau_{\text{merged}}$ 

```

887 A.2 Details of Reinforced Task Agents

888 In this section, we provide the detailed specifications
 889 and sources for the reinforced agentic models
 890 and base models used in our experiments. All models
 891 are publicly available on Hugging Face.

892 **Qwen2.5-7B-Instruction Series:** We utilize the
 893 following agents initialized from the Qwen2.5-7B-
 894 Instruct:

895 *Base Model (Qwen2.5-7B-Instruct)* (Yang et al.,
 896 2024)

897 😊 Qwen2.5-7B-Instruct

898 *Coding Agent (CURE)* (Wang et al., 2025b)

899 😊 ReasonFlux-Coder-7B

900 *Tool Agent (ToolRL)* (Qian et al., 2025)

901 😊 Qwen2.5-7B-Instruct-ToolRL-grpo-cold

902 *Memory Agent (MemAgent)* (Yu et al., 2025a)

903 😊 RL-MemoryAgent-7B

904 *Search Agent (ZeroSearch)* (Sun et al., 2025)

905 😊 ZeroSearch_google_V2_Qwen2.5_7B_Instruct

906 *AutoTIR Agent* (Wei et al., 2025a)

907 😊 AutoTIR-Qwen2.5-7B-Instruct

908 **Llama-3.2-3B-Instruction Series** To verify gen-
 909 eralization across architectures, we utilize the fol-
 910 lowing agents based on Llama-3.2-3B-Instruct:

911 *Base Model (Llama-3.2-3B-Instruct)* (Grattafiori
 912 et al., 2024)

913 😊 Llama-3.2-3B-Instruct

914 *Math Agent (GRPO-Math)* (Guo et al., 2025)

915 😊 Llama-3.2-3B-Instruct-GRPO-MATH-1EPOCH

916 *Tool Agent (ToolRL)* (Qian et al., 2025)

917 😊 ToolRL-Llama3.2-3B

918 *Search Agent (ZeroSearch)* (Sun et al., 2025)

919 😊 ZeroSearch_google_V2_Llama_3.2_3B_Instruct

920 B Evaluation Details

921 B.1 Coding Evaluation

922 Following the evaluation setting established in
 923 Wang et al. (2025b), we conduct a comprehensive
 924 evaluation of coding capabilities across five widely
 925 adopted coding benchmarks.

926 **Datasets.** We utilize **LiveBench** (White
 927 et al., 2025) (standard test set), **MBPP** (Austin
 928 et al., 2021) (standard test set), and **Live-
 929 CodeBench** (Jain et al., 2024) (Version 2, 511
 930 problems). For competition-level tasks, we include
 931 **CodeContests** (Li et al., 2022), filtering for
 932 tasks with difficulty ≤ 2 and utilizing a held-out
 933 split of 200 examples. Additionally, we use
 934 **CodeForces** (Penedo et al., 2025), comprising
 935 500 randomly sampled examples distinct from
 936 CodeContests.

937 **Evaluation Protocol.** To ensure consistency, all
 938 datasets are standardized to the stdio format. Func-
 939 tional inputs from LiveBench, LiveCodeBench, and
 940 MBPP are converted by placing variables on sep-
 941 arate lines and flattening lists. For verification,
 942 we use official ground-truth solutions for Code-
 943 Contests and MBPP. For the remaining datasets
 944 (CodeForces, LiveCodeBench, LiveBench), we uti-
 945 lize high-quality reference solutions generated by
 946 QwQ-32B (via Best-of-3 sampling). We employ
 947 vLLM (Kwon et al., 2023) for generation. Follow-
 948 ing standard practices, sampling parameters are set
 949 to temperature $T = 1.0$, top- $p = 0.95$. We report
 950 performance using pass accuracy, including code
 951 pass (ACC) and unit test pass (UT) (Pass@1) and
 952 Best-of-N (BoN, N=4) metrics.

953 B.2 Long-Context Memory Evaluation

954 To rigorously assess the long-context memory
955 capabilities of our agents, we adopt the evaluation
956 protocol from the RULER benchmark (Hsieh et al.,
957 2024), strictly following the data synthesis con-
958 figurations established in Yu et al. (2025a). We
959 specifically select RULER-HotpotQA and RULER-
960 SQuAD as our primary benchmarks to evaluate
961 multi-hop reasoning and precise fact retrieval.

962 **Datasets.** We utilize the following two tasks
963 adapted for long-context evaluation:

- 964 • **RULER-HotpotQA:** This task serves as a
965 robust testbed for *multi-hop reasoning*. In
966 this setup, multiple "golden paragraphs" con-
967 taining necessary evidence are embedded
968 within a vast amount of distractor content (the
969 haystack). The model must effectively iden-
970 tify and synthesize these scattered pieces of
971 evidence from its memory to correctly answer
972 complex questions.
- 973 • **RULER-SQuAD:** Adapted from the SQuAD
974 dataset, this task evaluates precise reading
975 comprehension. Ground-truth passages are
976 inserted into long distractor texts, requiring
977 the model to maintain high fidelity to specific
978 facts over extended sequences. This tests the
979 agent's ability to accurately recall specific in-
980 structions or details without hallucination.

981 **Evaluation Protocol.** Consistent with Yu et al.
982 (2025a), we synthesize test samples with varying
983 context lengths to stress-test memory capacity with
984 different context lengths (8K-128K for SQuAD
985 and 7K-896K for HotPotQA). The primary evalua-
986 tion metric is the Substring Exact Match (sub_em)
987 of the generated answers. High accuracy in this
988 setting demonstrates that the merged agent suc-
989 cessfully retains critical task-specific memory capabili-
990 ties and can effectively filter out noise (distractors)
991 inherent in long-context processing.

992 B.3 Tool Use Evaluation

993 To comprehensively evaluate the tool-use (func-
994 tion calling) capabilities of our agents, we em-
995 ploy the Berkeley Function Calling Leaderboard
996 (BFCL) (Patil et al., 2025), widely recognized as
997 the standard benchmark for assessing LLM agentic
998 behaviors. We specifically use the *Live* and *Non-*
999 *Live* datasets to measure performance across both
1000 real-world and synthetic scenarios.

1001 **Datasets.** We utilize the following two subsets to
1002 assess distinct dimensions of function calling:

- 1003 • **Non-Live Dataset (Synthetic & Curated):**
1004 Derived from BFCL V1, this subset con-
1005 sists of expert-curated synthetic tasks de-
1006 signed to test fundamental logic across var-
1007 ious languages (Python, Java, JavaScript)
1008 and SQL. It evaluates the model's adher-
1009 ence to precise instructions in controlled
1010 environments. The tasks of Non-Live
1011 datasets include: Multiple, Parallel,
1012 Relevance, Simple, Parallel_multiple
1013 and Irrelevance.
- 1014 • **Live Dataset (Real-World & Crowd-
1015 sourced):** Introduced in BFCL V2, this
1016 subset comprises user-contributed examples
1017 from real-world agent interactions. Un-
1018 like the Non-Live set, these samples are
1019 diverse and noisy, involving complex APIs
1020 with nested parameters. This benchmark
1021 specifically challenges the model's robust-
1022 ness in handling ambiguous queries and
1023 detecting function irrelevance, including
1024 seven tasks: Multiple, Irrelevance,
1025 Simple_java, Simple_javascript,
1026 Parallel_multiple, Parallel and
1027 Simple_python.

1028 **Evaluation Protocol.** To ensure robust evalua-
1029 tion, we utilize the Abstract Syntax Tree (AST)
1030 matching method provided by the BFCL frame-
1031 work. Unlike simple string matching, AST evalua-
1032 tion parses generated function calls into syntax
1033 trees to structurally verify argument permutations
1034 and formatting variations while enforcing strict
1035 type correctness. We report accuracy for both Live
1036 and Non-Live splits, with a particular focus on
1037 the challenging *Parallel* and *Parallel Multiple* cate-
1038 gories to demonstrate advanced planning capabili-
1039 ties.

1040 B.4 Evaluation Details for Llama-based 1041 Agents

1042 We evaluate agents trained from LLama3.2-3B-
1043 Instruction on three domains: math, search, and
1044 tool-use. For the math domain, we evaluate
1045 the model on GSM8K (Cobbe et al., 2021) and
1046 MATH500 (Hendrycks et al., 2021) datasets, pro-
1047 vided by LM-Evaluation-Harness (Gao et al., 2024).
1048 For the search domain, we follow the evaluation
1049 setting provided in ZeroSearch (Sun et al., 2025)

and evaluate the agent on NQ (Kwiatkowski et al., 2019) and 2WikiMultiHopQA (Ho et al., 2020). For tool-use, the setting is the same as Section 5.1. We take the average score across tasks for each domain for evaluation.

C Additional Experiments

C.1 Additional Reinforced Task Vector Analysis

To further verify heterogeneity in reinforced task vectors introduced by Section 3.2, we extend the number, domains, and architecture of the reinforced agents.

First, we include extra agents, web search agent ZeroSearch (Sun et al., 2025) and tool-integrated reasoning agent AutoTIR (Wei et al., 2025a), which are both RL-trained from Qwen2.5-Instruct-7B. Figure 7 shows that when including five agents specialized in multiple domains together, the heterogeneity in sparsity and distribution remains significant. Specifically, the sparsity of task vectors spans a wide spectrum, ranging from merely 3.2% for the Code agent to 54.3% for the Memory agent. The overlap analysis further reveals distinct behaviors: the Code agent is highly entangled with others, with 43.7% of its changed parameters shared among 3 or 4 other agents (w. 3-4). In contrast, agents like Tool and Memory maintain higher independence, with unique parameter ratios (w. 0) of 26.1% and 21.9%, respectively.

Second, we extend the experiment to Llama architecture and additional domains. We choose Llama3.2-3B-Instruction (Grattafiori et al., 2024) as the base model, and choose models trained from it via RL: search agent ZeroSearch (Sun et al., 2025), math reasoning agent (Zhao et al., 2025), and tool-using agent ToolRL (Qian et al., 2025). Figure 8 demonstrates that similar heterogeneity in task vectors persists across different model architectures. As shown in the left panel, the sparsity of task vectors varies significantly, ranging from 17.0% for the Math agent to 56.8% for the Tool agent. The overlap distribution (right panel) further highlights this diversity: the Tool agent modifies a large proportion of unique parameters (54.0%), whereas the Math agent shares the majority of its updates with other tasks, with only 24.2% of its modified parameters being unique. This confirms that the diverse characteristics of reinforced task vectors are consistent across different base models and task domains.

C.2 Instruction Following Evaluation

A primary concern in model merging, particularly when combining agents fine-tuned via RL on disparate domains, is the potential degradation of the base model’s general instruction following capabilities (i.e., catastrophic forgetting). To rigorously evaluate whether RAM compromises the model’s ability to follow general instructions while pursuing task specialization, we conducted evaluations on the **IFEval** (Instruction Following Evaluation) benchmark (Zhou et al., 2023) provided by LM-Evaluation-Harness (Gao et al., 2024). We report results across four metrics: Instruction Accuracy and Prompt Accuracy, under both Loose and Strict evaluation criteria. The results are presented in Table 3. We evaluated two sets of models:

- **Qwen2.5-7B-Instruct:** The primary setting used in the main paper, trained via RL to obtain the Coding, Tool, and Memory agents.
- **Llama-3.2-3B-Instruction:** To test the generalization of our method on a different architecture and scale, we use the fine-tuned Math, Tool, and Search agents based on Llama-3.2-3B-Instruction.

On the Qwen-based agents, RAM not only retains the general capabilities of the Base model but explicitly outperforms it across most metrics (e.g., +1.44 in Loose Instruction Accuracy and +1.56 in Strict Instruction Accuracy). This suggests that the specialized reasoning circuits preserved by RAM’s disjoint merging strategy can positively transfer to general instruction following. RAM+ shows a slight trade-off, generally maintaining parity with the base model on instruction-level metrics while incurring minor regressions in prompt-level accuracy. Merging on smaller Llama-based models is inherently more challenging due to limited parameter redundancy. While all merging methods exhibit some regression compared to the Base model, RAM demonstrates superior stability with less forgetting on instruction following. Notably, baseline methods like TIES and DARE+TIES suffer from severe performance collapse, dropping over 10 percentage points (e.g., -11.15 in Loose Instruction Accuracy). Even in strict evaluation, TIES fails to maintain robustness. In contrast, RAM avoids this collapse, showing significantly smaller regressions (approx. 2-3) and proving it is a much safer merging strategy for smaller architectures compared to aggressive trimming methods.




Figure 7: Additional distribution analysis for sparse task vectors.

Model	Qwen2.5-7B-Instruct				Llama-3.2-3B-Instruction			
	Instruction Acc (%)		Prompt Acc (%)		Instruction Acc (%)		Prompt Acc (%)	
	Loose	Strict	Loose	Strict	Loose	Strict	Loose	Strict
<i>Base and Task Experts</i>								
Base Model	69.18	64.39	58.96	53.23	69.90	63.31	59.15	51.20
Code/Math	73.14 (+3.96)	68.11 (+3.72)	62.48 (+3.52)	58.75 (+5.52)	68.59 (-1.31)	61.99 (-1.32)	57.49 (-1.66)	48.43 (-2.77)
Tool	71.82 (+2.64)	66.07 (+1.68)	61.55 (+2.59)	54.34 (+1.11)	68.11 (-1.79)	61.87 (-1.44)	57.12 (-2.03)	49.35 (-1.85)
Memory/Search	69.78 (+0.60)	66.07 (+1.68)	58.41 (-0.55)	53.60 (+0.37)	67.63 (-2.27)	60.67 (-2.64)	55.45 (-3.70)	46.58 (-4.62)
<i>Merged Models</i>								
Task Arithmetic (TA)	70.74 (+1.56)	65.71 (+1.32)	58.90 (-0.06)	53.23 (0.00)	68.82 (-1.08)	61.87 (-1.44)	57.12 (-2.03)	49.72 (-1.48)
Fisher	72.06 (+2.88)	66.91 (+2.52)	61.18 (+2.22)	55.08 (+1.85)	67.87 (-2.03)	61.75 (-1.56)	55.82 (-3.33)	49.17 (-2.03)
TIES	71.34 (+2.16)	67.39 (+3.00)	59.52 (+0.56)	55.64 (+2.41)	58.75 (-11.15)	58.87 (-4.44)	45.84 (-13.31)	46.21 (-4.99)
DARE+TA	70.38 (+1.20)	65.95 (+1.56)	58.41 (-0.55)	53.05 (-0.18)	59.11 (-10.79)	54.20 (-9.11)	46.21 (-12.94)	41.04 (-10.16)
DARE+TIES	69.42 (+0.24)	65.11 (+0.72)	57.67 (-1.29)	52.31 (-0.92)	58.63 (-11.27)	53.48 (-9.83)	45.66 (-13.49)	39.56 (-11.64)
RAM (Ours)	70.62 (+1.44)	65.95 (+1.56)	59.70 (+0.74)	53.23 (0.00)	67.39 (-2.51)	61.39 (-1.92)	55.45 (+3.70)	47.87 (-3.33)
RAM+ (Ours)	69.18 (0.00)	64.75 (+0.36)	57.86 (-1.10)	51.94 (-1.29)	66.19 (-3.71)	61.27 (-2.04)	53.97 (-5.18)	47.50 (-3.70)

Table 3: Evaluation of General Capabilities on IFEval benchmark. We report **Instruction Accuracy** and **Prompt Accuracy** under both **Loose** and **Strict** criteria. Values in parentheses denote the absolute change relative to the Base Model. Green values indicate improvement, while orange values indicate regression. RAM demonstrates superior robustness compared to TIES/DARE, especially on the smaller Llama-3.2-3B-Instruction model.

C.3 Merging Efficiency

Beyond merged performance, computational efficiency is another important factor for practical model merging. Figure 9 illustrates the comparison between merging time (in seconds) and the average score across benchmarks. Current baselines exhibit a clear dichotomy: methods like TA and Fisher are computationally efficient (<110s) but suffer from suboptimal performance (around 58.2), while complex methods like TIES and DARE variants achieve better scores at the cost of significant computational overhead (>400s). In contrast, our proposed methods occupy the Pareto frontier of the efficiency-performance landscape. Specifically, RAM achieves a remarkable score of 64.82 in just 75.4 seconds, surpassing DARE+TA in performance while offering a 5.5× speedup. Even our more intensive variant, RAM+, establishes a new SOTA performance (66.55) while remaining sig-

nificantly faster than both TIES and DARE. These results demonstrate that RAM effectively identifies critical parameters for merging processing without the extensive computational redundancy found in previous SOTAs.

C.4 Detailed Numerical Results for Pairwise Agent Merging

In Section 5.2, we visualize the performance of merging two agents using bar charts to highlight overall trends and comparative advantages. Here, we provide the corresponding detailed numerical results for all pairwise agent combinations, including **Coding+Tool**, **Tool+Memory**, and **Coding+Memory**, as shown in Tables 4, 5, and 6, respectively. These tables serve as a precise quantitative complement to the bar chart visualizations, enabling a fine-grained inspection of per-domain and per-metric behaviors.

1150
1151 Beyond merged performance, computational ef-
1152 ficiency is another important factor for practical
1153 model merging. Figure 9 illustrates the comparison
1154 between merging time (in seconds) and the aver-
1155 age score across benchmarks. Current baselines
1156 exhibit a clear dichotomy: methods like TA and
1157 Fisher are computationally efficient (<110s) but
1158 suffer from suboptimal performance (around 58.2),
1159 while complex methods like TIES and DARE vari-
1160 ants achieve better scores at the cost of signifi-
1161 cant computational overhead (>400s). In contrast,
1162 our proposed methods occupy the Pareto frontier
1163 of the efficiency-performance landscape. Specif-
1164 ically, RAM achieves a remarkable score of 64.82
1165 in just 75.4 seconds, surpassing DARE+TA in
1166 performance while offering a 5.5× speedup. Even
1167 our more intensive variant, RAM+, establishes a new
1168 SOTA performance (66.55) while remaining sig-
1169 nificantly faster than both TIES and DARE. These
1170 results demonstrate that RAM effectively identifies
1171 critical parameters for merging processing without
1172 the extensive computational redundancy found in
1173 previous SOTAs.
1174
1175

Model	Code				Tool				Avg
	LiveBench		LiveCodeBench		Live		Non-Live		
	ACC	UT	ACC	UT	Para	P_Mul	Para	P_Mul	
<i>Base and Task Models</i>									
Base	28.35	40.87	23.43	36.42	<u>56.25</u>	41.67	68.00	55.00	43.75
CURE (Coding)	37.70	49.27	30.23	45.76	<u>56.25</u>	37.50	64.00	51.50	46.53
ToolRL (Tool)	31.84	41.36	26.76	42.05	<u>56.25</u>	58.33	91.00	89.00	54.57
<i>Merged Models</i>									
TA	<u>37.11</u>	49.09	29.55	45.83	50.00	37.50	86.00	59.00	49.26
Fisher	36.72	50.29	30.68	48.18	50.00	37.50	88.50	64.50	50.80
TIES	35.74	46.50	<u>31.21</u>	44.48	<u>56.25</u>	54.17	83.00	<u>90.50</u>	55.23
DARE+TA	35.74	47.40	30.04	44.62	<u>56.25</u>	<u>58.33</u>	90.00	84.50	55.86
DARE+TIES	35.54	45.79	<u>31.21</u>	44.48	<u>56.25</u>	66.67	88.50	85.50	56.74
RAM	39.45	51.42	31.21	46.51	<u>56.25</u>	66.67	91.00	91.00	59.19
RAM+	39.45	<u>51.10</u>	32.53	47.55	62.50	66.67	90.50	90.00	60.04

Table 4: **Detailed results of model merging with code and tool using agents.** **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across all evaluated models, including the Task Models.

Model	Tool				Memory				Avg
	Live		Non-Live		HotpotQA		RulerQA		
	Para	P_Mul	Para	P_Mul	7k	14k	32k	64k	
<i>Base and Task Models</i>									
Base	<u>56.25</u>	41.67	68.00	55.00	60.94	50.00	64.84	58.59	56.91
ToolRL (Tool)	<u>56.25</u>	58.33	91.00	89.00	58.59	48.44	59.38	46.95	63.49
MemAgent (Memory)	37.50	50.00	78.50	48.50	78.91	78.12	81.25	77.34	66.27
<i>Merged Models</i>									
TA	<u>56.25</u>	62.50	91.00	90.00	72.66	70.31	<u>76.56</u>	74.22	74.19
Fisher	<u>56.25</u>	58.33	89.50	<u>91.00</u>	63.28	57.81	65.62	64.84	68.33
TIES	43.75	54.17	73.50	88.50	71.09	71.09	<u>76.56</u>	73.44	69.01
DARE+TA	62.50	58.33	89.00	90.00	57.03	55.47	63.28	59.38	66.87
DARE+TIES	62.50	58.33	90.00	91.50	77.34	71.97	75.00	<u>78.12</u>	75.60
RAM	56.25	66.67	90.00	89.50	76.56	78.12	78.12	78.12	76.67
RAM+	56.25	62.50	89.50	90.00	76.56	78.13	74.22	79.69	75.86

Table 5: **Detailed results of model merging with memory and tool using agents.** **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across all evaluated models, including the Task Models.

Model	Code				Memory				Avg
	LiveBench		LiveCodeBench		HotpotQA		RulerQA		
	ACC	UT	ACC	UT	7k	14k	32k	64k	
<i>Base and Task Models</i>									
Base	28.35	40.87	23.43	36.42	60.94	50.00	64.84	58.59	45.43
CURE (Coding)	37.70	49.27	30.23	45.76	58.59	<u>56.25</u>	60.94	44.22	47.87
MemAgent (Memory)	39.25	50.12	28.92	44.8	78.91	78.12	81.25	77.34	59.84
<i>Merged Models</i>									
TA	39.06	<u>52.25</u>	32.34	47.89	71.88	70.31	75.00	75.00	57.97
Fisher	38.67	50.53	31.46	46.58	57.03	50.78	58.59	55.47	48.64
TIES	38.87	50.88	32.42	<u>47.80</u>	74.22	71.09	78.91	72.66	58.36
DARE+TA	37.69	47.77	31.16	44.51	78.91	78.12	78.12	80.47	59.59
DARE+TIES	39.25	49.95	31.02	44.52	75.78	75.78	76.56	76.56	58.68
RAM	37.89	49.56	31.75	47.17	78.12	78.12	82.03	82.03	60.83
RAM+	41.21	53.64	31.16	46.46	78.91	76.56	79.69	82.03	61.21

Table 6: **Detailed results of model merging with coding and memory agents.** **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across all evaluated models, including the Task Models.

Model	Code				Tool				Memory				Avg
	LiveBench		LiveCodeBench		Live		Non-Live		HotpotQA		RulerQA		
	ACC	UT	ACC	UT	Para	P_Mul	Para	P_Mul	7k	14k	32k	64k	
RAM	38.28	49.71	31.96	47.72	56.25	66.67	91.00	91.50	75.78	75.00	74.22	79.69	64.82
RAM+	40.23	52.57	31.60	46.84	56.25	70.83	90.50	91.00	79.69	78.13	78.91	82.03	66.55
RAM+s	39.84	51.83	30.82	46.59	56.25	70.83	90.00	91.50	77.34	76.56	76.56	77.34	65.46

Table 7: The comparison results of a soft-saturation rescaling transformation.




Figure 8: The task vector analysis for agents trained from Llama3.2-3B-Instruction via RL. **Left:** Sparsity of task vectors varies between tool-using agent, web search agent, and the math reasoning agent models. **Right:** The number of overlaps with other task vectors.




Figure 9: RAM/RAM+ demonstrates a superior trade-off between merging time and average score compared to the baseline method.

Coding + Tool. As reported in Table 6, RAM and RAM+ consistently achieve the highest average performance among all merged models. In particular, RAM+ attains an average score of 60.04, significantly outperforming the strongest baseline DARE+TIES (56.74). Notably, RAM/RAM+ improve both coding accuracy (LiveBench / LiveCodeBench) and complex tool-use metrics (Parallel and Parallel-Multiple), indicating that preserving task-unique reinforced updates enables the merged model to simultaneously retain algorithmic reasoning and structured tool invocation capabilities. In contrast, baseline methods exhibit a clear trade-off, improving one domain at the expense of the other due to signal dilution in sparse RL task vectors.

Tool + Memory. Table 5 presents the detailed results for merging Tool and Memory agents. RAM achieves the highest overall average score (76.67), while RAM+ remains highly competitive (75.86), both surpassing all baselines. RAM-based models show strong performance on long-context memory tasks (HotpotQA and RulerQA) without degrading tool-use accuracy, particularly on challenging Parallel and Parallel-Multiple subsets. These results demonstrate that RAM effectively isolates memory-specific reinforced updates from tool-specific ones, preventing destructive interference that commonly

occurs in averaging-based merging methods.

Coding + Memory. As shown in Table 6, merging Coding and Memory agents further stresses the heterogeneity of reinforced task vectors, as these two domains exhibit minimal parameter overlap. Despite this challenge, RAM+ achieves the best average performance (61.21), outperforming all baselines and even exceeding the original Coding agent on LiveBench ACC/UT. This highlights the effectiveness of overlap-aware rescaling in compensating for performance loss in shared subspaces, while fully preserving unique reasoning and memory patterns critical for each task.

Across all pairwise combinations, the numerical results reported here are fully consistent with the trends observed in the bar charts in the main text. Specifically, RAM and RAM+ not only deliver the highest average scores but also exhibit superior robustness across heterogeneous domains and metrics. These findings further confirm that the advantages of RAM are not limited to tri-agent merging, but generalize naturally to arbitrary agent combinations, reinforcing its suitability as a unified merging framework for reinforced agentic models.

1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1211
1212
1213

1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237

1238
1239

C.5 Additional Results in Merging Three Agents

1240
1241
1242
1243

In Section 5.2, we present the agent performance in three domains. Here, we further provide experiment results on additional tasks and settings to comprehensively verify the advantages of RAM.

1244
1245
1246
1247
1248
1249
1250
1251

Overall Analysis. Tables 8, 9, and 10 report comprehensive results of merging three reinforced agents across coding, tool-use, and long-context memory domains, substantially extending the representative results in Section 5.2. These results consistently corroborate the advantages of RAM and RAM+ under a wide range of evaluation metrics and task granularities.

1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270

Coding Domain. As shown in Table 8, RAM-based methods achieve the strongest overall performance among merged models across LiveBench, LiveCodeBench, MBPP, and CodeContests. Notably, RAM attains the highest average score (49.64), outperforming all baseline merging strategies, while RAM+ further improves performance on challenging subsets such as LiveBench ACC/UT and MBPP ACC/UT. Importantly, RAM and RAM+ frequently match or exceed the original Coding agent on multiple metrics (highlighted in red), indicating that preserving and selectively amplifying task-unique reinforced updates effectively avoids signal dilution that hampers prior methods. In contrast, methods relying on global averaging or random dropping (TA, Fisher, DARE) exhibit inconsistent gains and fail to simultaneously retain high unit-test robustness and generalization accuracy.

1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286

Tool-Use Domain. Table 9 presents detailed tool-use evaluations over both Live and Non-Live settings. Across diverse function-calling scenarios—including parallel, parallel-multiple, irrelevance detection, and language-specific tasks—RAM achieves the best average score (77.51), while RAM+ remains a close second (77.45), both surpassing all existing merging baselines. Crucially, RAM-based models consistently excel in structurally complex settings such as Live Parallel_multiple and Non-Live Parallel, where signal dilution in sparse RL task vectors is most detrimental. These results demonstrate that RAM effectively preserves tool-specific reasoning circuits while still benefiting from shared knowledge introduced by other agents.

1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318

Long-Context Memory Domain. As reported in Table 10, RAM and RAM+ show clear dominance on long-context memory benchmarks across all context lengths. RAM+ achieves the highest overall average score (77.57), outperforming both merged baselines and, in several settings, the specialized MemAgent itself. In particular, RAM+ consistently delivers state-of-the-art performance on long-context HotpotQA (112K–896K) and Ruler-SQuAD (16K–64K), confirming that overlap-aware rescaling effectively compensates for performance degradation introduced by averaging shared subspaces. By contrast, Fisher and Task Arithmetic suffer from substantial performance drops at long context lengths, reflecting their inability to preserve sparse, task-specific memory-related updates.

1309
1310
1311
1312
1313
1314
1315
1316
1317
1318

Across all three domains, the additional results reinforce three central conclusions. First, reinforced task vectors exhibit strong heterogeneity, making uniform merging strategies fundamentally suboptimal. Second, preserving and rescaling task-unique parameters is essential for maintaining expert-level performance after merging. Third, RAM and RAM+ consistently outperform prior SOTA merging methods not only in average performance but also in robustness across metrics, datasets, and context scales. Together, these findings provide strong empirical evidence that RAM is a principled and effective solution for merging multiple RL-trained agents into a unified generalist model.

D Baseline Details

D.1 Baselines with Signal Dilution

1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336

Here we provide a detailed introduction of baselines and explain signal dilution happen in them. In summary, despite their distinct algorithmic designs, Task Arithmetic, TIES-Merging, and DARE all inevitably succumb to Signal Dilution in the RL setting due to a shared inability to distinguish between *shared consensus* and *unique specialization*. Specifically, Task Arithmetic is forced to apply a conservative global scaling ($\lambda \approx 1/N$) to prevent magnitude explosion in shared subspaces, collaterally suppressing unique task vectors that require no scaling. Similarly, TIES-Merging fails to strictly isolate disjoint parameters due to “noise drift” in inactive RL models, which erroneously inflates the normalization denominator with meaningless artifacts. Even DARE, despite introducing a rescaling

Model	LiveBench				LiveCodeBench				MBPP				CodeContests				Avg
	ACC	UT	BoN ACC	BoN UT	ACC	UT	BoN ACC	BoN UT	ACC	UT	BoN ACC	BoN UT	ACC	UT	BoN ACC	BoN UT	
<i>Base and Task Models</i>																	
Base	28.35	40.87	30.07	48.56	23.43	36.42	26.67	42.26	62.72	68.50	70.91	76.84	18.20	28.35	23.01	32.06	41.08
CURE	37.70	49.27	46.01	58.96	30.23	45.76	38.36	55.60	70.36	76.38	79.64	85.78	26.05	36.81	32.21	45.61	50.92
ToolRL	31.84	41.36	32.81	43.78	26.76	42.05	30.14	46.05	66.97	72.31	76.02	81.29	22.18	32.94	26.78	37.71	44.44
MemAgent	39.25	50.12	39.84	48.48	28.92	44.80	32.16	48.63	67.63	73.99	71.04	77.84	21.86	31.68	24.69	36.09	46.06
<i>Merged Models</i>																	
TA	38.09	51.62	43.75	52.97	31.95	46.69	35.61	46.44	69.68	<u>75.71</u>	74.21	80.39	25.73	36.70	28.45	38.70	48.54
Fisher	36.72	48.73	45.31	57.79	30.87	45.89	35.81	53.26	68.67	74.92	75.57	83.23	<u>24.26</u>	<u>35.87</u>	27.20	39.38	48.97
TIES	<u>39.25</u>	49.88	44.53	53.08	30.63	46.32	36.59	<u>54.30</u>	67.30	73.35	74.66	82.34	26.05	<u>37.00</u>	30.50	42.05	49.24
DARE+TA	37.50	48.60	<u>46.10</u>	56.41	31.94	46.69	37.18	53.44	69.68	<u>75.71</u>	<u>78.73</u>	83.83	23.74	32.82	<u>28.45</u>	38.02	49.30
DARE+TIES	35.93	45.66	46.87	57.88	29.26	39.53	38.55	51.76	<u>70.70</u>	76.34	80.64	86.22	21.13	27.62	27.20	35.93	48.20
RAM	38.28	49.71	42.97	57.98	<u>31.96</u>	47.72	<u>37.45</u>	54.36	69.46	75.56	76.47	81.89	24.06	35.07	26.78	44.51	49.64
RAM+	40.23	52.57	46.10	57.20	31.60	46.84	34.64	51.56	70.93	76.98	77.38	84.13	22.80	33.21	23.01	35.46	49.04

Table 8: **Additional results of model merging on coding domains.** **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across *all evaluated models*, including the Task Models.

Models	Live				Non-live				Avg					
	Multiple	Parallel	Relevance	Simple	Parallel_multiple	Irrelevance	Multiple	Irrelevance	S_java	S_javascript	Parallel_multiple	Parallel	S_python	
<i>Base and Task Models</i>														
Base	58.59	56.25	87.50	68.99	41.67	68.21	77.50	77.50	54.00	62.00	55.00	68.00	87.25	66.34
CURE	59.54	56.25	81.25	69.77	37.50	68.33	76.50	77.92	58.00	60.00	51.50	64.00	91.50	65.54
ToolRL	76.83	56.25	93.75	79.07	58.33	71.95	94.00	82.92	62.00	62.00	89.00	91.00	93.50	77.74
MemAgent	73.98	37.50	93.75	69.77	50.00	56.11	82.50	71.67	66.00	66.00	48.50	78.50	90.25	68.04
<i>Merged Models</i>														
TA	76.16	43.75	93.75	75.19	45.83	68.55	91.00	79.17	60.00	62.00	69.50	87.50	94.75	72.86
Fisher	74.55	43.75	93.75	74.03	41.67	70.02	91.50	81.25	58.00	58.00	63.50	86.50	94.50	71.62
TIES	<u>75.50</u>	43.75	87.50	75.19	54.17	63.46	92.50	<u>76.25</u>	65.00	<u>66.00</u>	67.50	84.00	94.50	72.71
DARE+TA	<u>75.50</u>	50.00	93.75	76.36	<u>62.50</u>	63.57	92.00	74.17	65.00	66.00	89.50	92.50	93.75	76.50
DARE+TIES	75.31	56.25	93.75	76.36	58.33	63.80	<u>92.00</u>	76.67	65.00	68.00	90.00	<u>91.50</u>	93.25	76.94
RAM	<u>75.50</u>	56.25	93.75	75.58	66.67	67.99	92.00	76.67	65.00	64.00	91.50	91.00	91.75	77.51
RAM+	75.12	56.25	93.75	75.58	66.67	67.87	92.00	76.67	65.00	66.00	<u>90.50</u>	90.00	91.50	77.45

Table 9: **Additional results of model merging on tool using.** **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across *all evaluated models*, including the Task Models.

Model	Ruler_SQuAD												Ruler_HotpotQA				Avg
	8K	16K	32K	64K	128K	7K	14K	28K	56K	112K	224K	448K	896K				
<i>Base and Task Models</i>																	
Base	65.63	62.50	64.84	58.59	56.25	60.94	50.00	51.56	48.44	42.19	36.72	27.34	25.78	50.06			
CURE	46.80	61.72	60.94	44.22	61.72	58.59	56.25	46.88	47.66	40.63	35.93	42.97	35.16	49.19			
ToolRL	62.50	57.81	59.38	46.95	67.18	58.59	48.44	51.56	45.31	42.97	42.19	35.16	35.94	50.31			
MemAgent	83.59	78.12	81.25	77.34	81.25	78.91	78.12	81.25	77.34	79.69	72.66	76.56	72.66	78.36			
<i>Merged Models</i>																	
TA	72.66	70.31	73.44	72.66	72.66	69.53	68.75	68.75	67.19	63.28	64.84	57.03	48.44	66.89			
Fisher	60.16	67.97	58.59	60.94	67.97	60.06	49.22	49.22	49.22	39.06	32.81	35.94	35.94	51.32			
TIES	71.88	75.00	75.00	75.78	73.44	71.88	82.03	71.97	68.75	67.97	70.31	62.50	64.06	71.58			
DARE+TA	71.88	77.34	<u>77.34</u>	70.31	80.47	<u>76.56</u>	76.56	72.65	71.88	<u>70.31</u>	73.44	71.88	<u>70.31</u>	73.92			
DARE+TIES	75.78	78.13	76.56	74.22	78.91	75.00	<u>77.34</u>	75.00	70.36	70.43	76.56	74.44	74.22	75.15			
RAM	79.68	82.03	74.22	79.69	76.56	75.78	75.00	78.91	75.78	70.31	75.78	71.09	74.22	76.08			
RAM+	77.34	82.81	78.91	82.03	79.68	79.69	78.13	78.91	75.56	69.53	78.13	73.43	74.22	77.57			

Table 10: **Additional results of model merging on memory for long-context tasks.** **Bold** and underlined values denote the best and second-best performance among *merged models*, respectively. Cells highlighted in red indicate the best performance across *all evaluated models*, including the Task Models.

mechanism, only compensates for *internal* dropout rather than *external* merging, thereby remaining dependent on the global scaling of Task Arithmetic to function. Consequently, all three paradigms effectively drive the magnitude of task-specific updates towards $\frac{1}{N}\tau$, underscoring the necessity of a topology-aware merging framework.

Fisher Fisher merging (Matena and Raffel, 2022) improves upon uniform averaging by weighing parameters according to their diagonal Fisher information F , which approximates the posterior precision (or local curvature) of the model. Formally, the merged update is computed as $\tau_{\text{merged}} = \frac{\sum_i F_i \tau_i}{\sum_i F_i}$. However, this method remains susceptible to signal dilution in the sparse RL setting. Consider a "unique" parameter updated solely by task t (where $\tau_t \neq 0$ and $\tau_{i \neq t} = 0$). Crucially, the inactive models ($i \neq t$) still contribute to the normalization term in the denominator, as their Fisher values F_i —representing the confidence of the pre-trained base model—are typically non-zero. Consequently, the effective scaling factor for the unique signal becomes $\frac{F_t}{F_t + \sum_{i \neq t} F_i}$. Since the denominator accumulates the "inertia" (precision) of all inactive models, the task-specific update τ_t is inevitably scaled down, mirroring the dilution effect observed in uniform averaging as the number of tasks N increases.

Task Arithmetic. Task Arithmetic (Ilharco et al., 2023) constructs a multi-task model by linearly combining task vectors, typically expressed as $\tau_{\text{merged}} = \sum_i \lambda \tau_i$. While effective for disjoint tasks in SFT, it faces a critical trade-off in the RL setting due to the heterogeneity of parameter overlap. To prevent catastrophic magnitude shifts in *shared* subspaces (where multiple agents push parameters in similar directions), the scaling factor λ must be conservative (typically $\lambda \approx 1/N$) to maintain the merged weights within a valid optimization landscape. However, this global scaling creates a structural conflict: *unique* task vectors, which reside in non-overlapping subspaces and do not suffer from additive interference, are subjected to the same aggressive downscaling. Consequently, the update for a task-specific parameter is reduced to $\frac{1}{N}\tau_t$, effectively suppressing the critical, idiosyncratic behaviors required for expert-level performance in domains like tool-use or memory.

TIES-Merging. TIES-Merging (Yadav et al., 2023) attempts to mitigate interference by keep-

ing only the top- $k\%$ magnitude parameters (Trimming) and calculating a disjoint mean. However, its reliance on a *uniform* trimming rate k creates a critical vulnerability when handling the **heterogeneous sparsity** of RL task vectors. As shown in Figure 2, RL agents exhibit vastly different update densities (e.g., Code: $\sim 3\%$, Memory: $\sim 54\%$). Applying a fixed k (e.g., 20%) inevitably leads to a dilemma: it either over-trims dense vectors (losing info) or, more disastrously, under-trims sparse vectors. For a sparse agent like the Code model, a standard k retains a large volume of noise parameters alongside the true signal. These noise parameters, mistakenly treated as valid updates, overlap with the active parameters of other tasks, inflating the normalization factor in the disjoint mean calculation ($\tau_m = \frac{1}{|A_p|} \sum \tau_t$). Consequently, the critical, unique signal from one task is averaged with noise from others, resulting in signal dilution.

DARE. DARE (Yu et al., 2024) randomly drops parameters and enlarges others to reduce redundancy, theoretically approximating the original task vector's expectation. However, DARE fails to address signal dilution for two reasons. First, its factor $(1/(1-p))$ is designed solely to compensate for the internal dropout rate p , not for the external averaging with other models. When combined with Task Arithmetic (DARE+TA or DARE+TIES), the global merging scale λ must still be kept small (e.g., $1/N$) to stabilize shared parameters, inevitably downscaling the unique, sparsity-preserved updates. Second, DARE assumes that parameter updates are highly redundant (a property of SFT), whereas RL updates are inherently sparse and functionally essential. Randomly dropping parameters in an already sparse RL vector risks severing critical reasoning circuits, which cannot be recovered simply by rescaling the remaining weights.

D.2 Hyperparameters Search

Table 11 shows the searched ranges of model merging methods' hyperparameters. We search for the best performance for evaluation and comparison.

E Rescaling Variant

To bridge the theoretical requirement with numerical stability mentioned in Section 4.2, besides clipping, we further explore a soft-saturation transformation $\phi(x) = \frac{x}{1+x}$ to map the unbounded ratio ρ_t to the bounded interval $[0, 1]$. Our operational

1337
1338
1339
1340
1341
1342
1343

1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405

1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364

1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424

1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386

1425
1426
1427
1428
1429
1430
1431
1432
1433
1434

Table 11: Searched ranges of hyperparameters of model merging baselines.

Model Merging Methods	Search Ranges of Hyperparameters
Task Arithmetic	scaling term to merge model parameters: [0.1, 0.3, 0.5, 0.7, 0.9, 1.0]
Fisher	scaling term to merge model parameters: [0.1, 0.3, 0.5, 0.7, 0.9, 1.0], number of examples to compute Fisher information matrix: [256, 512, 1024, 2048]
TIES	scaling term to merge model parameters: [0.1, 0.3, 0.5, 0.7, 0.9, 1.0], ratio to retain parameters with largest-magnitude values: [0.1, 0.2, 0.3]
DARE	search the drop rate p in [0.1, 0.2, ..., 0.9]

1435 scaling rule is thus defined as:

$$1436 \quad \lambda_t = 1 + r \cdot \left(\frac{\rho_t}{1 + \rho_t} \right). \quad (8)$$

1437 Here, the hyperparameter r absorbs the dilution
1438 factor $(1 - \beta)$.

1439 It is worth noting that this normalized formulation
1440 is mathematically equivalent to the ratio of
1441 shared parameters to the *total* active parameters.
1442 By substituting the definition of ρ_t from Eq. 2, we
1443 obtain:

$$1444 \quad \frac{\rho_t}{1 + \rho_t} = \frac{\sum_{i:c_i \geq 2} m_{t,i}}{\|\mathbf{m}_t\|_0}. \quad (9)$$

1445 Here, the numerator sums the shared parameters,
1446 while the denominator $\|\mathbf{m}_t\|_0$ represents the total
1447 count of active parameters (satisfying $\|\mathbf{m}_t\|_0 =$
1448 $\sum_{i:c_i \geq 2} m_{t,i} + \sum_{i:c_i=1} m_{t,i}$). This transformation
1449 provides a dual advantage: it retains the monotonicity
1450 derived from the conservation principle (higher
1451 overlap yields higher compensation) while enforc-
1452 ing a strict upper bound to ensure optimization
1453 robustness. Table 7 presents the comparative per-
1454 formance of the soft-saturation rescaling strategy,
1455 denoted as **RAM+s**. Empirically, RAM+s achieves
1456 an average score of 65.46, consistently outper-
1457 forming the non-rescaled baseline RAM (64.82)
1458 across all three domains. This reinforces our core
1459 hypothesis that compensating for signal dilution
1460 in unique parameter subspaces is essential for re-
1461 covering expert capabilities, regardless of the spe-
1462 cific scaling function used. However, RAM+s
1463 slightly underperforms compared to the clipped
1464 linear variant (RAM+, 66.55). While the soft-
1465 saturation transformation $\phi(x) = \frac{x}{1+x}$ offers a
1466 theoretically elegant, strictly bounded mapping,
1467 it appears to dampen the scaling factor more ag-
1468 gressively than the linear approach. This conser-
1469 vatism limits performance in tasks requiring robust
1470 signal preservation, such as Long-Context Mem-
1471 ory, where RAM+s scores 77.34 on RulerQA (64k)

1472 compared to RAM+'s 82.03. Conversely, in the
1473 Tool domain, RAM+s remains highly effective,
1474 matching RAM+ with a score of 70.83 on Live
1475 Parallel-Multiple tasks. These findings suggest that
1476 while the soft-saturation rule provides a stable alter-
1477 native, the clipped linear rule offers a superior
1478 trade-off between signal amplification and numeri-
1479 cal stability.

1472
1473
1474
1475
1476
1477
1478
1479